

CS315: DATABASE SYSTEMS

RELATIONAL ALGEBRA

Arnab Bhattacharya

`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,
Indian Institute of Technology, Kanpur

<http://web.cse.iitk.ac.in/~cs315/>

2nd semester, 2018-19

Mon 12:00-13:15, Tue 9:00-10:15

Relational Algebra

- *Procedural language* to specify database queries
- Operators are functions from one or two input relations to an output relation
 - 1 Select: σ
 - 2 Project: Π
 - 3 Union: \cup
 - 4 Set Difference: $-$
 - 5 Cartesian Product: \times
 - 6 Rename: ρ

Relational Algebra

- *Procedural language* to specify database queries
- Operators are functions from one or two input relations to an output relation
 - 1 Select: σ
 - 2 Project: Π
 - 3 Union: \cup
 - 4 Set Difference: $-$
 - 5 Cartesian Product: \times
 - 6 Rename: ρ
- Uses *propositional calculus* consisting of *expressions* connected by
 - 1 and: \wedge
 - 2 or: \vee
 - 3 not: \neg
- Each term is of the form
 $\langle \text{attr/const} \rangle \text{ comparator } \langle \text{attr/const} \rangle$
where comparator is one of $=, \neq, >, \geq, <, \leq$

Select

- $\sigma_p(r) = \{t | t \in r \text{ and } p(t)\}$
- p is called the **selection predicate**
- Select all tuples from r that satisfies the predicate p
- Schema is

Select

- $\sigma_p(r) = \{t | t \in r \text{ and } p(t)\}$
- p is called the **selection predicate**
- Select all tuples from r that satisfies the predicate p
- Schema is not changed
- Applying $\sigma_{A=B \wedge D > 5}$ on

A	B	C	D
1	1	2	7
1	2	5	7
2	2	9	3
2	2	8	6

returns

Select

- $\sigma_p(r) = \{t | t \in r \text{ and } p(t)\}$
- p is called the **selection predicate**
- Select all tuples from r that satisfies the predicate p
- Schema is not changed
- Applying $\sigma_{A=B \wedge D > 5}$ on

A	B	C	D
1	1	2	7
1	2	5	7
2	2	9	3
2	2	8	6

returns	A	B	C	D
	1	1	2	7
	2	2	8	6

- $\Pi_{A_1, \dots, A_k}(r)$
- A_i , etc. are attributes of r
- Select only the specified attributes A_1, \dots, A_k from all tuples of r
- Duplicate rows are removed, since relations are sets

- $\Pi_{A_1, \dots, A_k}(r)$
- A_i , etc. are attributes of r
- Select only the specified attributes A_1, \dots, A_k from all tuples of r
- Duplicate rows are removed, since relations are sets
- Schema is

- $\Pi_{A_1, \dots, A_k}(r)$
- A_i , etc. are attributes of r
- Select only the specified attributes A_1, \dots, A_k from all tuples of r
- Duplicate rows are removed, since relations are sets
- Schema is changed
- Applying $\Pi_{A,C}$ on

A	B	C
1	1	5
1	2	5
2	3	5
2	4	8

returns

Project

- $\Pi_{A_1, \dots, A_k}(r)$
- A_i , etc. are attributes of r
- Select only the specified attributes A_1, \dots, A_k from all tuples of r
- Duplicate rows are removed, since relations are sets
- Schema is changed
- Applying $\Pi_{A,C}$ on

			A	B	C
			<hr/>		
			1	1	5
			1	2	5
			2	3	5
			2	4	8
			<hr/>		
				A	C
returns				1	5
				2	5
				2	8

Set Union

- $r \cup s = \{t | t \in r \text{ or } t \in s\}$
- Relations r and s must have the same *arity* (i.e., number of attributes)
- They must have same *type* of attribute in each column as well, i.e., attribute domains must be *compatible*
- If attribute names are not same, renaming should be used
- Schema is

Set Union

- $r \cup s = \{t | t \in r \text{ or } t \in s\}$
- Relations r and s must have the same *arity* (i.e., number of attributes)
- They must have same *type* of attribute in each column as well, i.e., attribute domains must be *compatible*
- If attribute names are not same, renaming should be used
- Schema is not changed
- Applying \cup on

A	B		A	B
1	1	and	1	2
1	2		2	3
2	1			

returns

Set Union

- $r \cup s = \{t \mid t \in r \text{ or } t \in s\}$
- Relations r and s must have the same *arity* (i.e., number of attributes)
- They must have same *type* of attribute in each column as well, i.e., attribute domains must be *compatible*
- If attribute names are not same, renaming should be used
- Schema is not changed
- Applying \cup on

A	B	and	A	B
1	1		1	2
1	2		2	3
2	1			

returns	A	B
	1	1
	1	2
	2	1
	2	3

Set Difference

- $r - s = \{t | t \in r \text{ and } t \notin s\}$
- Relations r and s must have the same *arity* (i.e., number of attributes)
- They must have same *type* of attribute in each column as well, i.e., attribute domains must be *compatible*
- If attribute names are not same, renaming should be used
- Schema is

Set Difference

- $r - s = \{t | t \in r \text{ and } t \notin s\}$
- Relations r and s must have the same *arity* (i.e., number of attributes)
- They must have same *type* of attribute in each column as well, i.e., attribute domains must be *compatible*
- If attribute names are not same, renaming should be used
- Schema is not changed
- Applying – on

A	B		A	B
1	1	and	1	2
1	2		2	3
2	1			

returns

Set Difference

- $r - s = \{t | t \in r \text{ and } t \notin s\}$
- Relations r and s must have the same *arity* (i.e., number of attributes)
- They must have same *type* of attribute in each column as well, i.e., attribute domains must be *compatible*
- If attribute names are not same, renaming should be used
- Schema is not changed
- Applying – on

A	B		A	B
1	1		1	2
1	2	and	2	3
2	1			

	A	B
returns	1	1
	2	1

Cartesian Product

- $r \times s = \{t \mid t \in r \text{ and } t \in s\}$
- Attributes of relations r and s should be disjoint
- If attributes are not disjoint, renaming should be used
- Schema is

Cartesian Product

- $r \times s = \{t \mid t \in r \text{ and } t \in s\}$
- Attributes of relations r and s should be disjoint
- If attributes are not disjoint, renaming should be used
- Schema is changed
- Applying \times on

A	B		C	D	E
1	1	and	1	2	7
2	2		2	6	8
			5	7	9

returns

Cartesian Product

- $r \times s = \{t \mid t \in r \text{ and } t \in s\}$
- Attributes of relations r and s should be disjoint
- If attributes are not disjoint, renaming should be used
- Schema is changed
- Applying \times on

<table><tr><th>A</th><th>B</th></tr><tr><td>1</td><td>1</td></tr><tr><td>2</td><td>2</td></tr></table>		A	B	1	1	2	2	and	<table><tr><th>C</th><th>D</th><th>E</th></tr><tr><td>1</td><td>2</td><td>7</td></tr><tr><td>2</td><td>6</td><td>8</td></tr><tr><td>5</td><td>7</td><td>9</td></tr></table>			C	D	E	1	2	7	2	6	8	5	7	9
A	B																						
1	1																						
2	2																						
C	D	E																					
1	2	7																					
2	6	8																					
5	7	9																					
returns		A	B	C	D	E																	
		1	1	1	2	7																	
		1	1	2	6	8																	
		1	1	5	7	9																	
		2	2	1	2	7																	
		2	2	2	6	8																	
	2	2	5	7	9																		

Rename

- $\rho_N(E)$ returns E , but under the new name N
- For n -ary relations, $\rho_{N(A_1, \dots, A_n)}(E)$ returns the result of expression E , but under the new name N and the attributes renamed to A_1 , etc.
- Schema is

Rename

- $\rho_N(E)$ returns E , but under the new name N
- For n -ary relations, $\rho_{N(A_1, \dots, A_n)}(E)$ returns the result of expression E , but under the new name N and the attributes renamed to A_1 , etc.
- Schema is changed although its meaning is not
- Applying $\rho_{S(C,D)}$ on $r(A, B)$

A	B
1	1
1	2
2	3
2	4

returns

Rename

- $\rho_N(E)$ returns E , but under the new name N
- For n -ary relations, $\rho_{N(A_1, \dots, A_n)}(E)$ returns the result of expression E , but under the new name N and the attributes renamed to A_1 , etc.
- Schema is changed although its meaning is not
- Applying $\rho_{S(C,D)}$ on $r(A, B)$

		A	B
		<hr/>	
		1	1
		1	2
		2	3
		2	4
		C	D
		<hr/>	
		1	1
returns		1	2
		2	3
		2	4

Additional Operations

- Additional operators have been defined
 - 1 Set Intersection: \cap
 - 2 Join: \bowtie
 - 3 Division: \div
 - 4 Assignment: \leftarrow
- These do not add any power to the basic relational algebra
 - They can be defined using the six basic operators
- However, they simplify queries

Set Intersection

- $r \cap s = \{t | t \in r \text{ and } t \in s\}$
- Relations r and s must have the same *arity* (i.e., number of attributes)
- They must have same *type* of attribute in each column as well, i.e., attribute domains must be *compatible*
- If attribute names are not same, renaming should be used
- Schema is

Set Intersection

- $r \cap s = \{t \mid t \in r \text{ and } t \in s\}$
- Relations r and s must have the same *arity* (i.e., number of attributes)
- They must have same *type* of attribute in each column as well, i.e., attribute domains must be *compatible*
- If attribute names are not same, renaming should be used
- Schema is not changed
- Applying \cap on

A	B		A	B
1	1	and	1	2
1	2		2	3
2	1			

returns

Set Intersection

- $r \cap s = \{t \mid t \in r \text{ and } t \in s\}$
- Relations r and s must have the same *arity* (i.e., number of attributes)
- They must have same *type* of attribute in each column as well, i.e., attribute domains must be *compatible*
- If attribute names are not same, renaming should be used
- Schema is not changed
- Applying \cap on

A	B		A	B
1	1		1	2
1	2	and	2	3
2	1			
returns				
A	B			
1	2			

Set Intersection

- $r \cap s = \{t | t \in r \text{ and } t \in s\}$
- Relations r and s must have the same *arity* (i.e., number of attributes)
- They must have same *type* of attribute in each column as well, i.e., attribute domains must be *compatible*
- If attribute names are not same, renaming should be used
- Schema is not changed
- Applying \cap on

A	B		A	B
1	1	and	1	2
1	2		2	3
2	1			

returns

A	B
1	2

- $r \cap s = r - (r - s)$

Join

- $r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$
- Join is too common a query to not have its own operator
- Schema is

Join

- $r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$
- Join is too common a query to not have its own operator
- Schema is same as $r \times s$ but (potentially) less number of tuples
- The above form is the most general, called the **theta join**

Join

- $r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$
- Join is too common a query to not have its own operator
- Schema is same as $r \times s$ but (potentially) less number of tuples
- The above form is the most general, called the **theta join**
- **Equality join**: When the join condition only contains equality
 - $r \bowtie_{B=C} s$

- $r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$
- Join is too common a query to not have its own operator
- Schema is same as $r \times s$ but (potentially) less number of tuples
- The above form is the most general, called the **theta join**
- **Equality join**: When the join condition only contains equality
 - $r \bowtie_{B=C} s$
- **Natural join**: If two relations share an attribute (also its *name*), equality join on that common attribute
 - Denoted by $*$ or simply \bowtie without any predicate
 - Changes schema by retaining only one copy of common attribute
 - $r * s = r \bowtie s = r \bowtie_{r.A=s.A} s$
- Applying \bowtie on

A	B		A	C	
1	1	and	1	2	returns
1	2		2	3	
2	1				

- $r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$
- Join is too common a query to not have its own operator
- Schema is same as $r \times s$ but (potentially) less number of tuples
- The above form is the most general, called the **theta join**
- **Equality join**: When the join condition only contains equality
 - $r \bowtie_{B=C} s$
- **Natural join**: If two relations share an attribute (also its *name*), equality join on that common attribute
 - Denoted by $*$ or simply \bowtie without any predicate
 - Changes schema by retaining only one copy of common attribute
 - $r * s = r \bowtie s = r \bowtie_{r.A=s.A} s$
- Applying \bowtie on

A	B		A	C		A	B	C
1	1	and	1	2	returns	1	1	2
1	2		2	3		1	2	2
2	1					2	1	3

Division

- $r \div s = \{t \mid t \in \Pi_{R-S}(r) \text{ and } \forall u \in s (tu \in r)\}$
- $q = r \div s$ is the largest relation satisfying $q \times s \subseteq r$

Division

- $r \div s = \{t \mid t \in \Pi_{R-S}(r) \text{ and } \forall u \in s (tu \in r)\}$
- $q = r \div s$ is the largest relation satisfying $q \times s \subseteq r$
- Relation r must have a schema that is a proper superset of the schema of s , i.e., $S \subset R$
- Used for queries of the form “for all”
- Schema is

Division

- $r \div s = \{t \mid t \in \Pi_{R-S}(r) \text{ and } \forall u \in s (tu \in r)\}$
- $q = r \div s$ is the largest relation satisfying $q \times s \subseteq r$
- Relation r must have a schema that is a proper superset of the schema of s , i.e., $S \subset R$
- Used for queries of the form “for all”
- Schema is changed to $R - S$
- Applying \div on

A	B	
1	5	
1	6	
1	7	
2	5	and $\frac{B}{5}$ returns
2	6	6
3	5	
3	7	
4	5	

Division

- $r \div s = \{t \mid t \in \Pi_{R-S}(r) \text{ and } \forall u \in s (tu \in r)\}$
- $q = r \div s$ is the largest relation satisfying $q \times s \subseteq r$
- Relation r must have a schema that is a proper superset of the schema of s , i.e., $S \subset R$
- Used for queries of the form “for all”
- Schema is changed to $R - S$
- Applying \div on

A	B					
1	5					
1	6					
1	7					
2	5	and	<table><tr><th>B</th></tr><tr><td>5</td></tr><tr><td>6</td></tr></table>	B	5	6
B						
5						
6						
2	6		returns			
3	5		<table><tr><th>A</th></tr><tr><td>1</td></tr><tr><td>2</td></tr></table>	A	1	2
A						
1						
2						
3	7					
4	5					

Division (contd.)

- Applying \div on

A	B	C	D		C	D	
1	5	2	7		2	7	returns
1	5	3	7		3	7	
1	6	3	7	and			
2	6	2	7				
2	6	3	7				
3	6	2	7				
3	6	3	7				
3	5	3	7				

Division (contd.)

- Applying \div on

A	B	C	D
1	5	2	7
1	5	3	7
1	6	3	7
2	6	2	7
2	6	3	7
3	6	2	7
3	6	3	7
3	5	3	7

and

C	D
2	7
3	7

 returns

A	B
1	5
2	6
3	6

Division (contd.)

- Applying \div on

A	B	C	D		C	D		A	B
1	5	2	7					1	5
1	5	3	7					2	6
1	6	3	7					3	6
2	6	2	7	and	2	7	returns		
2	6	3	7		3	7			
3	6	2	7						
3	6	3	7						
3	5	3	7						

- $r \div s = \Pi_{R-S}(r) - \Pi_{R-S}((\Pi_{R-S}(r) \times s) - \Pi_{R-S,S}(r))$

Assignment

- $s \leftarrow E(r)$ assigns the relation resulting from applying E on r to s
- Useful in complex queries to hold intermediate values
 - Can be used sequentially
- Schema is

Assignment

- $s \leftarrow E(r)$ assigns the relation resulting from applying E on r to s
- Useful in complex queries to hold intermediate values
 - Can be used sequentially
- Schema is not changed
- $\rho_{(A=B)}(\Pi_{A,B}(r))$ can be broken into $s \leftarrow \Pi_{A,B}(r)$ and $\rho_{(A=B)}(s)$

r		
A	B	C
1	1	7
2	2	8
5	7	9

Assignment

- $s \leftarrow E(r)$ assigns the relation resulting from applying E on r to s
- Useful in complex queries to hold intermediate values
 - Can be used sequentially
- Schema is not changed
- $\rho_{(A=B)}(\Pi_{A,B}(r))$ can be broken into $s \leftarrow \Pi_{A,B}(r)$ and $\rho_{(A=B)}(s)$

r		
A	B	C
1	1	7
2	2	8
5	7	9

s	
A	B
1	1
2	2
5	7

and

$\rho_{(A=B)}(s)$	
A	B
1	1
2	2

Composition of Operators

- Expressions can be built using multiple operators
- Applying $\sigma_{A=C}(r \times s)$ on

A	B	and	C	D	E	intermediately produces
1	1		1	2	7	
2	2		2	6	8	
			5	7	9	

A	B	C	D	E	and finally returns	A	B	C	D	E
1	1	1	2	7		1	1	1	2	7
1	1	2	6	8		2	2	2	6	8
1	1	5	7	9						
2	2	1	2	7						
2	2	2	6	8						
2	2	5	7	9						

Precedence and Associativity

- Precedence is generally assumed to be

Precedence	Operators
Highest	σ, Π, ρ
Medium	$\bowtie, \bowtie_{\theta}, \times$
Lowest	$\cup, \cap, -$

- Associativity is assumed to be left-to-right
- Not part of definition
- Therefore, best to use explicit brackets

Example Schema

- course (code, title, *ctype*, webpage)
- coursetype (ctype, *dept*)
- faculty (fid, name, *dept*, designation)
- department (deptid, name)
- semester (yr, half)
- offering (*coursecode*, *yr*, *half*, *instructor*)
- student (roll, name, *dept*)
- program (*roll*, *pctype*)
- registration (*coursecode*, *roll*, *yr*, *half*, grade)

Example Queries

- Find all courses offered in the year 2018

Example Queries

- Find all courses offered in the year 2018
 $\sigma_{yr=2018}(\text{offering})$

Example Queries

- Find all courses offered in the year 2018

$\sigma_{yr=2018}(\text{offering})$

- Find the course codes for all courses offered in the year 2018

Example Queries

- Find all courses offered in the year 2018

$\sigma_{yr=2018}(\text{offering})$

- Find the course codes for all courses offered in the year 2018

$\Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$

Example Queries

- Find all courses offered in the year 2018

$\sigma_{yr=2018}(\text{offering})$

- Find the course codes for all courses offered in the year 2018

$\Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$

- Find the course codes for all the courses offered in either of the years 2017 and 2018

Example Queries

- Find all courses offered in the year 2018

$$\sigma_{yr=2018}(\text{offering})$$

- Find the course codes for all courses offered in the year 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

- Find the course codes for all the courses offered in either of the years 2017 and 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2017}(\text{offering})) \cup \Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

Example Queries

- Find all courses offered in the year 2018

$$\sigma_{yr=2018}(\text{offering})$$

- Find the course codes for all courses offered in the year 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

- Find the course codes for all the courses offered in either of the years 2017 and 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2017}(\text{offering})) \cup \Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

- Find the course codes for all the courses offered in the year 2017 but not in 2018

Example Queries

- Find all courses offered in the year 2018

$$\sigma_{yr=2018}(\text{offering})$$

- Find the course codes for all courses offered in the year 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

- Find the course codes for all the courses offered in either of the years 2017 and 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2017}(\text{offering})) \cup \Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

- Find the course codes for all the courses offered in the year 2017 but not in 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2017}(\text{offering})) - \Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

Example Queries

- Find all courses offered in the year 2018

$$\sigma_{yr=2018}(\text{offering})$$

- Find the course codes for all courses offered in the year 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

- Find the course codes for all the courses offered in either of the years 2017 and 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2017}(\text{offering})) \cup \Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

- Find the course codes for all the courses offered in the year 2017 but not in 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2017}(\text{offering})) - \Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

- Find the course codes for all the courses offered in both the years 2017 and 2018

Example Queries

- Find all courses offered in the year 2018

$$\sigma_{yr=2018}(\text{offering})$$

- Find the course codes for all courses offered in the year 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

- Find the course codes for all the courses offered in either of the years 2017 and 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2017}(\text{offering})) \cup \Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

- Find the course codes for all the courses offered in the year 2017 but not in 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2017}(\text{offering})) - \Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

- Find the course codes for all the courses offered in both the years 2017 and 2018

$$\Pi_{\text{coursecode}}(\sigma_{yr=2017}(\text{offering})) \cap \Pi_{\text{coursecode}}(\sigma_{yr=2018}(\text{offering}))$$

Example Queries

- Find the titles of all courses offered in the year 2018

Example Queries

- Find the titles of all courses offered in the year 2018

$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\sigma_{\text{offering.coursecode}=\text{courses.coursecode}}(\text{offering} \times \text{courses})))$

Example Queries

- Find the titles of all courses offered in the year 2018

$$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\sigma_{\text{offering.coursecode}=\text{courses.coursecode}}(\text{offering} \times \text{courses})))$$
$$\Pi_{\text{title}}(\sigma_{\text{offering.coursecode}=\text{courses.coursecode}}(\sigma_{\text{yr}=2018}(\text{offering}) \times \text{courses}))$$

Example Queries

- Find the titles of all courses offered in the year 2018

$$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\sigma_{\text{offering.coursecode}=\text{courses.coursecode}}(\text{offering} \times \text{courses})))$$
$$\Pi_{\text{title}}(\sigma_{\text{offering.coursecode}=\text{courses.coursecode}}(\sigma_{\text{yr}=2018}(\text{offering}) \times \text{courses}))$$
$$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\text{offering} \bowtie \text{courses}))$$

Example Queries

- Find the titles of all courses offered in the year 2018

$$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\sigma_{\text{offering.coursecode}=\text{courses.coursecode}}(\text{offering} \times \text{courses})))$$

$$\Pi_{\text{title}}(\sigma_{\text{offering.coursecode}=\text{courses.coursecode}}(\sigma_{\text{yr}=2018}(\text{offering}) \times \text{courses}))$$

$$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\text{offering} \bowtie \text{courses}))$$

$$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\text{offering}) \bowtie \text{courses})$$

Example Queries

- Find the titles of all courses offered in the year 2018

$$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\sigma_{\text{offering.coursecode}=\text{courses.coursecode}}(\text{offering} \times \text{courses})))$$

$$\Pi_{\text{title}}(\sigma_{\text{offering.coursecode}=\text{courses.coursecode}}(\sigma_{\text{yr}=2018}(\text{offering}) \times \text{courses}))$$

$$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\text{offering} \bowtie \text{courses}))$$

$$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\text{offering}) \bowtie \text{courses})$$

- Find the years when *all* the courses of type 5 were offered

Example Queries

- Find the titles of all courses offered in the year 2018

$$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\sigma_{\text{offering.coursecode}=\text{courses.coursecode}}(\text{offering} \times \text{courses})))$$

$$\Pi_{\text{title}}(\sigma_{\text{offering.coursecode}=\text{courses.coursecode}}(\sigma_{\text{yr}=2018}(\text{offering}) \times \text{courses}))$$

$$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\text{offering} \bowtie \text{courses}))$$

$$\Pi_{\text{title}}(\sigma_{\text{yr}=2018}(\text{offering}) \bowtie \text{courses})$$

- Find the years when *all* the courses of type 5 were offered

$$\text{ct} \leftarrow \rho_{\text{coursecode}}(\Pi_{\text{code}}(\sigma_{\text{ctype}=5}(\text{course})))$$

$$(\Pi_{\text{coursecode}, \text{yr}}(\text{offering})) \div \text{ct}$$

Extended Relational Algebra

- The power of relational algebra can be enhanced by
 - 1 Generalized Projection
 - 2 Grouping and Aggregate operations
 - 3 Outer Join

Generalized Projection

- Extends project operator by allowing arbitrary arithmetic functions in attribute list
- $\Pi_{F_1, \dots, F_k}(E)$
- F_i , etc. are arithmetic expressions involving constants and attributes in schema of E
- Applying $\Pi_{B-A, 2C}$ on r

A	B	C
1	1	5
1	2	5
2	3	5
2	4	8

returns

Generalized Projection

- Extends project operator by allowing arbitrary arithmetic functions in attribute list
- $\Pi_{F_1, \dots, F_k}(E)$
- F_i , etc. are arithmetic expressions involving constants and attributes in schema of E
- Applying $\Pi_{B-A, 2C}$ on r

	A	B	C
	1	1	5
	1	2	5
	2	3	5
	2	4	8
returns	B-A		2C
	0		10
	1		10
	2		16

Aggregate Operations

- Aggregate functions that can be used are *avg*, *min*, *max*, *sum*, *count*
- Can be applied on groups of tuples as well
- Aggregate operation is of the form $G_1, \dots, G_k \mathcal{G}_{F_1(A_1), \dots, F_n(A_n)}(E)$ where
 - G_1, \dots, G_k is the list of attributes on which to group (may be empty)
 - Each F_i is an aggregate function that operates on the attribute A_i
- Applying $\mathcal{G}_{sum(C)}$ on r

A	B	C	
1	1	5	
1	2	5	returns
2	3	5	
2	4	8	

Aggregate Operations

- Aggregate functions that can be used are *avg*, *min*, *max*, *sum*, *count*
- Can be applied on groups of tuples as well
- Aggregate operation is of the form $G_1, \dots, G_k \mathcal{G}_{F_1(A_1), \dots, F_n(A_n)}(E)$ where
 - G_1, \dots, G_k is the list of attributes on which to group (may be empty)
 - Each F_i is an aggregate function that operates on the attribute A_i
- Applying $\mathcal{G}_{sum(C)}$ on r

A	B	C	returns $\frac{sum(C)}{23}$
1	1	5	
1	2	5	
2	3	5	
2	4	8	

Aggregate Operations

- First, the tuples are grouped according to G_1, \dots, G_k
- Then, aggregate functions $F_1(A_1), \dots, F_n(A_n)$ are applied on each group
- Schema changes to $(G_1, \dots, G_k, F_1(A_1), \dots, F_n(A_n))$
- Applying $AG_{sum(C)}$ on r

A	B	C	returns
1	1	5	
1	2	5	
2	3	5	
2	4	8	
3	4	8	

Aggregate Operations

- First, the tuples are grouped according to G_1, \dots, G_k
- Then, aggregate functions $F_1(A_1), \dots, F_n(A_n)$ are applied on each group
- Schema changes to $(G_1, \dots, G_k, F_1(A_1), \dots, F_n(A_n))$
- Applying $AG_{sum(C)}$ on r

A	B	C		A	sum(C)
1	1	5	returns	1	10
1	2	5		2	13
2	3	5		3	8
2	4	8			
3	4	8			

Aggregate Operations

- First, the tuples are grouped according to G_1, \dots, G_k
- Then, aggregate functions $F_1(A_1), \dots, F_n(A_n)$ are applied on each group
- Schema changes to $(G_1, \dots, G_k, F_1(A_1), \dots, F_n(A_n))$
- Applying $AG_{sum(C)}$ on r

A	B	C
1	1	5
1	2	5
2	3	5
2	4	8
3	4	8

returns

A	sum(C)
1	10
2	13
3	8

- Applying $AG_{sum(C)}$ on r

A	B	C
1	1	5
1	2	5
1	2	4

returns

Aggregate Operations

- First, the tuples are grouped according to G_1, \dots, G_k
- Then, aggregate functions $F_1(A_1), \dots, F_n(A_n)$ are applied on each group
- Schema changes to $(G_1, \dots, G_k, F_1(A_1), \dots, F_n(A_n))$
- Applying $AG_{sum(C)}$ on r

A	B	C
1	1	5
1	2	5
2	3	5
2	4	8
3	4	8

returns

A	sum(C)
1	10
2	13
3	8

- Applying $AG_{sum(C)}$ on r

A	B	C
1	1	5
1	2	5
1	2	4

returns

A	B	sum(C)
1	1	5
1	2	9

Outer Join

- Extension of the join to retain more information
- Computes join and then adds tuples to result that do not match
- Requires use of *null* values
- **Left outer join** $r \bowtie_{\theta} s$ retains *every* tuple from left or first relation
 - If no matching tuple is found in right or second relation, values are padded with *null*
- **Right outer join** $r \ltimes_{\theta} s$ is defined analogously
- **Full outer join** $r \Join_{\theta} s$ retains all tuples from both relations
 - Non-matching fields are filled with *null* values

Outer Join

- Extension of the join to retain more information
- Computes join and then adds tuples to result that do not match
- Requires use of *null* values
- **Left outer join** $r \bowtie_{\theta} s$ retains *every* tuple from left or first relation
 - If no matching tuple is found in right or second relation, values are padded with *null*
- **Right outer join** $r \ltimes_{\theta} s$ is defined analogously
- **Full outer join** $r \bowtie_{\theta} s$ retains all tuples from both relations
 - Non-matching fields are filled with *null* values
- Consequently, ordinary join is sometimes called **inner join**
- “Outer” word is sometimes dropped from join yielding **left join**, **right join** and **full join**

Outer Join

- Extension of the join to retain more information
- Computes join and then adds tuples to result that do not match
- Requires use of *null* values
- **Left outer join** $r \bowtie_{\theta} s$ retains *every* tuple from left or first relation
 - If no matching tuple is found in right or second relation, values are padded with *null*
- **Right outer join** $r \ltimes_{\theta} s$ is defined analogously
- **Full outer join** $r \bowtie_{\theta} s$ retains all tuples from both relations
 - Non-matching fields are filled with *null* values
- Consequently, ordinary join is sometimes called **inner join**
- “Outer” word is sometimes dropped from join yielding **left join**, **right join** and **full join**
- When no θ condition is specified, it is **natural outer join**

Outer Join Examples

A	B		A	C	
1	5		1	7	
2	6	\bowtie	2	8	$=$
3	7		4	9	

Outer Join Examples

A	B		A	C		A	B	C
1	5		1	7		1	5	7
2	6	\bowtie	2	8	=	2	6	8
3	7		4	9				

A	B		A	C	
1	5		1	7	
2	6	\bowtie	2	8	=
3	7		4	9	

Outer Join Examples

A	B		A	C		A	B	C
1	5	\bowtie	1	7	$=$	1	5	7
2	6		2	8		2	6	8
3	7		4	9				

A	B		A	C		A	B	C
1	5	\bowtie_{right}	1	7	$=$	1	5	7
2	6		2	8		2	6	8
3	7		4	9		3	7	null

A	B		A	C	
1	5	\bowtie_{left}	1	7	$=$
2	6		2	8	
3	7		4	9	

Outer Join Examples

A	B		A	C		A	B	C
1	5	\bowtie	1	7	$=$	1	5	7
2	6		2	8		2	6	8
3	7		4	9				

A	B		A	C		A	B	C
1	5	\bowtie_{right}	1	7	$=$	1	5	7
2	6		2	8		2	6	8
3	7		4	9		3	7	null

A	B		A	C		A	B	C
1	5	\bowtie_{left}	1	7	$=$	1	5	7
2	6		2	8		2	6	8
3	7		4	9		4	null	9

Outer Join Examples

A	B		A	C	
1	5		1	7	
2	6	\bowtie	2	8	$=$
3	7		4	9	

Outer Join Examples

A	B		A	C		A	B	C
1	5		1	7		1	5	7
2	6	\bowtie	2	8	=	2	6	8
3	7		4	9				

A	B		A	C	
1	5		1	7	
2	6	\bowtie	2	8	=
3	7		4	9	

Outer Join Examples

A	B		A	C		A	B	C
1	5		1	7		1	5	7
2	6	\bowtie	2	8	=	2	6	8
3	7		4	9				

A	B		A	C		A	B	C
1	5		1	7		1	5	7
2	6	\bowtie	2	8	=	2	6	8
3	7		4	9		3	7	null
						4	null	9

Example Queries

- Find the total number of courses offered in the year 2018

Example Queries

- Find the total number of courses offered in the year 2018

$$\mathcal{G}_{count(coursecode)}(\sigma_{yr=2018}(offering))$$

Example Queries

- Find the total number of courses offered in the year 2018

$$\mathcal{G}_{count(coursecode)}(\sigma_{yr=2018}(offering))$$

- For each instructor, find the total number of courses offered by her in the year 2018

Example Queries

- Find the total number of courses offered in the year 2018

$$\mathcal{G}_{count(coursecode)}(\sigma_{yr=2018}(offering))$$

- For each instructor, find the total number of courses offered by her in the year 2018

$$instructor \mathcal{G}_{count(coursecode)}(\sigma_{yr=2018}(offering))$$

Example Queries

- Find the total number of courses offered in the year 2018

$$\mathcal{G}_{count(coursecode)}(\sigma_{yr=2018}(offering))$$

- For each instructor, find the total number of courses offered by her in the year 2018

$$instructor \mathcal{G}_{count(coursecode)}(\sigma_{yr=2018}(offering))$$

- For each instructor, find the total number of courses offered by her per year

Example Queries

- Find the total number of courses offered in the year 2018

$$\mathcal{G}_{count(coursecode)}(\sigma_{yr=2018}(offering))$$

- For each instructor, find the total number of courses offered by her in the year 2018

$$instructor \mathcal{G}_{count(coursecode)}(\sigma_{yr=2018}(offering))$$

- For each instructor, find the total number of courses offered by her per year

$$instructor, yr \mathcal{G}_{count(coursecode)}(offering)$$

Example Queries

- Find the total number of courses offered in the year 2018

$$\mathcal{G}_{count(coursecode)}(\sigma_{yr=2018}(offering))$$

- For each instructor, find the total number of courses offered by her in the year 2018

$$instructor \mathcal{G}_{count(coursecode)}(\sigma_{yr=2018}(offering))$$

- For each instructor, find the total number of courses offered by her per year

$$instructor, yr \mathcal{G}_{count(coursecode)}(offering)$$

- For each course, indicate the most recent year it was offered

Example Queries

- Find the total number of courses offered in the year 2018

$$\mathcal{G}_{count(coursecode)}(\sigma_{yr=2018}(offering))$$

- For each instructor, find the total number of courses offered by her in the year 2018

$$instructor \mathcal{G}_{count(coursecode)}(\sigma_{yr=2018}(offering))$$

- For each instructor, find the total number of courses offered by her per year

$$instructor, yr \mathcal{G}_{count(coursecode)}(offering)$$

- For each course, indicate the most recent year it was offered

$$course-offering \leftarrow coursecode \mathcal{G}_{max(yr)}(offering)$$

$$course-year \leftarrow \rho_{(code, yr)}(\Pi_{coursecode, max(yr)}(course-offering))$$

$$course \bowtie course-year$$

Null Values

- Null denotes an unknown or missing value
- Arithmetic expressions involving null evaluate to null
- Aggregate functions ignore null
- Duplicate elimination and grouping treats null as any other value, i.e., two null values are same
 - $\text{null} = \text{null}$ evaluates to true

Truth Tables with Null Values

- Comparison with null otherwise returns *unknown*, not false
- If false is used, consider two expressions $\text{not}(A < 5)$ and $A \geq 5$ and when attribute contains null
 - They will not be the same
- Three-valued logic with *unknown*
 - Or
 - unknown or true = true
 - unknown or false = unknown
 - unknown or unknown = unknown
 - And
 - unknown and true = unknown
 - unknown and false = false
 - unknown and unknown = unknown
 - Not
 - not unknown = unknown
- Select operation treats unknown as false

Database Modification

- Contents of a database may be modified by
 - 1 Deletion
 - 2 Insertion
 - 3 Updating
- Assignment operator is used to express these operations

Deletion

- $r \leftarrow r - E$ deletes tuples in the result set of the query E from the relation r
- Only whole tuples can be deleted, not some attributes
- Applying $r \leftarrow r - \sigma_{A=1}(r)$ on

A	B	C
1	1	5
1	2	5
2	3	5
2	4	8

returns

Deletion

- $r \leftarrow r - E$ deletes tuples in the result set of the query E from the relation r
- Only whole tuples can be deleted, not some attributes
- Applying $r \leftarrow r - \sigma_{A=1}(r)$ on

A	B	C		A	B	C
1	1	5		2	3	5
1	2	5	returns	2	4	8
2	3	5				
2	4	8				

Insertion

- $r \leftarrow r \cup E$ inserts tuples in the result set of the query E into the relation r
- Only whole tuples can be inserted, not some attributes
- If a specific tuple needs to be inserted, E is specified as a relation containing only that tuple
- Applying $r \leftarrow r \cup \{(1, 2, 5)\}$ on

A	B	C
1	1	5
2	3	5
2	4	8

returns

Insertion

- $r \leftarrow r \cup E$ inserts tuples in the result set of the query E into the relation r
- Only whole tuples can be inserted, not some attributes
- If a specific tuple needs to be inserted, E is specified as a relation containing only that tuple
- Applying $r \leftarrow r \cup \{(1, 2, 5)\}$ on

A	B	C		A	B	C
1	1	5		1	1	5
2	3	5	returns	2	3	5
2	4	8		2	4	8
				1	2	5

Updation

- Updates allow values of only some attributes to change
- $r \leftarrow \Pi_{F_1, \dots, F_n}(r)$ where each F_i is
 - Either the i th attribute of r if it is not to be changed
 - Or the result of the expression F_i involving constants and attributes resulting in the new value of the i th attribute
- Applying $r \leftarrow \Pi_{A, 2*B, C}(r)$ on

A	B	C
1	2	5
1	1	5
2	4	8

returns

Updation

- Updates allow values of only some attributes to change
- $r \leftarrow \Pi_{F_1, \dots, F_n}(r)$ where each F_i is
 - Either the i th attribute of r if it is not to be changed
 - Or the result of the expression F_i involving constants and attributes resulting in the new value of the i th attribute
- Applying $r \leftarrow \Pi_{A, 2*B, C}(r)$ on

A	B	C		A	B	C
1	2	5	returns	1	4	5
1	1	5		1	2	5
2	4	8		2	8	8

- Applying $r \leftarrow \Pi_{A, 2*B, C}(\sigma_{A=1}(r))$ on

A	B	C	
1	2	5	returns
1	1	5	
2	4	8	

Updation

- Updates allow values of only some attributes to change
- $r \leftarrow \Pi_{F_1, \dots, F_n}(r)$ where each F_i is
 - Either the i th attribute of r if it is not to be changed
 - Or the result of the expression F_i involving constants and attributes resulting in the new value of the i th attribute
- Applying $r \leftarrow \Pi_{A, 2*B, C}(r)$ on

A	B	C		A	B	C
1	2	5	returns	1	4	5
1	1	5		1	2	5
2	4	8		2	8	8

- Applying $r \leftarrow \Pi_{A, 2*B, C}(\sigma_{A=1}(r))$ on

A	B	C		A	B	C
1	2	5	returns	1	4	5
1	1	5		1	2	5
2	4	8				

Example Queries

- Create a new department “Astronomy” with id 18

Example Queries

- Create a new department “Astronomy” with id 18

$\text{department} \leftarrow \text{department} \cup \{18, \text{“Astronomy”}\}$

Example Queries

- Create a new department “Astronomy” with id 18

$\text{department} \leftarrow \text{department} \cup \{18, \text{“Astronomy”}\}$

- Delete the course whose code is “CS200”

Example Queries

- Create a new department “Astronomy” with id 18

$\text{department} \leftarrow \text{department} \cup \{18, \text{“Astronomy”}\}$

- Delete the course whose code is “CS200”

$\text{course} \leftarrow \text{course} - \sigma_{\text{code}=\text{“CS200”}}(\text{course})$

Example Queries

- Create a new department “Astronomy” with id 18

$\text{department} \leftarrow \text{department} \cup \{18, \text{“Astronomy”}\}$

- Delete the course whose code is “CS200”

$\text{course} \leftarrow \text{course} - \sigma_{\text{code}=\text{“CS200”}}(\text{course})$

- Update the title of the course “DBMS” to “Database Systems”

Example Queries

- Create a new department “Astronomy” with id 18

$$\text{department} \leftarrow \text{department} \cup \{18, \text{“Astronomy”}\}$$

- Delete the course whose code is “CS200”

$$\text{course} \leftarrow \text{course} - \sigma_{\text{code}=\text{“CS200”}}(\text{course})$$

- Update the title of the course “DBMS” to “Database Systems”

$$\text{old-course} \leftarrow \sigma_{\text{title}=\text{“DBMS”}}(\text{course})$$
$$\text{updated-course} \leftarrow \Pi_{\text{code}, \text{“Database Systems”}, \text{ctype}, \text{webpage}}(\text{old-course})$$
$$\text{course} \leftarrow (\text{course} \cup \text{updated-course}) - \text{old-course}$$

Integrity Constraints Violations

- Deletion may violate

Integrity Constraints Violations

- Deletion may violate
 - **Referential integrity**: If a primary key is deleted, the corresponding foreign referencing key becomes orphan
 - Should be restricted (rejected) or cascaded or set to null
- Insertion may violate

Integrity Constraints Violations

- Deletion may violate
 - **Referential integrity**: If a primary key is deleted, the corresponding foreign referencing key becomes orphan
 - Should be restricted (rejected) or cascaded or set to null
- Insertion may violate
 - **Referential integrity**: If a foreign key is inserted, the corresponding primary referenced key must be present
 - Should be restricted

Integrity Constraints Violations

- Deletion may violate
 - **Referential integrity**: If a primary key is deleted, the corresponding foreign referencing key becomes orphan
 - Should be restricted (rejected) or cascaded or set to null
- Insertion may violate
 - **Referential integrity**: If a foreign key is inserted, the corresponding primary referenced key must be present
 - Should be restricted
 - **Domain constraint**: Value is outside the domain
 - Should be restricted or domain updated

Integrity Constraints Violations

- Deletion may violate
 - **Referential integrity**: If a primary key is deleted, the corresponding foreign referencing key becomes orphan
 - Should be restricted (rejected) or cascaded or set to null
- Insertion may violate
 - **Referential integrity**: If a foreign key is inserted, the corresponding primary referenced key must be present
 - Should be restricted
 - **Domain constraint**: Value is outside the domain
 - Should be restricted or domain updated
 - **Key constraint**: If an insertion violates the property of being a key
 - Should be restricted or design modified

Integrity Constraints Violations

- Deletion may violate
 - **Referential integrity**: If a primary key is deleted, the corresponding foreign referencing key becomes orphan
 - Should be restricted (rejected) or cascaded or set to null
- Insertion may violate
 - **Referential integrity**: If a foreign key is inserted, the corresponding primary referenced key must be present
 - Should be restricted
 - **Domain constraint**: Value is outside the domain
 - Should be restricted or domain updated
 - **Key constraint**: If an insertion violates the property of being a key
 - Should be restricted or design modified
 - **Entity integrity**: If the primary key of the inserted tuple is null
 - Should be restricted
- Updation may violate

Integrity Constraints Violations

- Deletion may violate
 - **Referential integrity**: If a primary key is deleted, the corresponding foreign referencing key becomes orphan
 - Should be restricted (rejected) or cascaded or set to null
- Insertion may violate
 - **Referential integrity**: If a foreign key is inserted, the corresponding primary referenced key must be present
 - Should be restricted
 - **Domain constraint**: Value is outside the domain
 - Should be restricted or domain updated
 - **Key constraint**: If an insertion violates the property of being a key
 - Should be restricted or design modified
 - **Entity integrity**: If the primary key of the inserted tuple is null
 - Should be restricted
- Updation may violate
 - *Referential integrity*
 - *Domain constraint*
 - *Key constraint*
 - *Entity integrity*

Drawbacks of Relational Algebra

Drawbacks of Relational Algebra

- First-order propositional logic
- Do not support recursive closure operations
 - Find supervisors of A at *all* levels
- Needs specifying multiple queries, each solving only one level at a time

Multiset Variant

- Relations are **multisets** or **bags** of tuples
 - Duplicate tuples are allowed
 - Select, project, set operations change

Multiset Variant

- Relations are **multisets** or **bags** of tuples
 - Duplicate tuples are allowed
 - Select, project, set operations change
- Distinct or duplicate elimination operator: δ
 - Removes duplicate tuples
 - Reduces relation to sets