

# CS315: DATABASE SYSTEMS SCHEDULES

Arnab Bhattacharya

`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,  
Indian Institute of Technology, Kanpur

<http://web.cse.iitk.ac.in/~cs315/>

2<sup>nd</sup> semester, 2018-19

Mon 12:00-13:15, Tue 9:00-10:15

# Schedule

- A **schedule** is a *chronological* sequence of instructions from *concurrent* transactions

# Schedule

- A **schedule** is a *chronological* sequence of instructions from *concurrent* transactions
- If a transaction appears in a schedule, *all* its instructions must appear in the schedule
- *Order* of instructions within a transaction must be maintained in the schedule

# Schedule

- A **schedule** is a *chronological* sequence of instructions from *concurrent* transactions
- If a transaction appears in a schedule, *all* its instructions must appear in the schedule
- *Order* of instructions within a transaction must be maintained in the schedule
- A transaction finishing successfully will have *commit* as the last instruction
- A transaction not finishing successfully will have *abort* as the last instruction
- Commit and abort statements may be omitted if obvious

# Example

- T1 transfers 50 from A to B and then T2 transfers 10% of A to B
- A **serial** schedule:  
 $r_1(A); A := A - 50; w_1(A); r_1(B); B := B + 50; w_1(B);$   
 $r_2(A); t := 0.1A; A := A - t; w_2(A); r_2(B); B := B + t; w_2(B);$

# Example

- T1 transfers 50 from A to B and then T2 transfers 10% of A to B
- A **serial** schedule:  
 $r_1(A); A := A - 50; w_1(A); r_1(B); B := B + 50; w_1(B);$   
 $r_2(A); t := 0.1A; A := A - t; w_2(A); r_2(B); B := B + t; w_2(B);$
- Another schedule:  
 $r_1(A); A := A - 50; w_1(A); r_2(A); t := 0.1A; A := A - t;$   
 $w_2(A); r_1(B); B := B + 50; w_1(B); r_2(B); B := B + t; w_2(B);$

# Example

- T1 transfers 50 from A to B and then T2 transfers 10% of A to B
- A **serial** schedule:  
 $r_1(A); A := A - 50; w_1(A); r_1(B); B := B + 50; w_1(B);$   
 $r_2(A); t := 0.1A; A := A - t; w_2(A); r_2(B); B := B + t; w_2(B);$
- Another schedule:  
 $r_1(A); A := A - 50; w_1(A); r_2(A); t := 0.1A; A := A - t;$   
 $w_2(A); r_1(B); B := B + 50; w_1(B); r_2(B); B := B + t; w_2(B);$
- This is not a serial schedule but is **equivalent**

# Example

- T1 transfers 50 from A to B and then T2 transfers 10% of A to B
- A **serial** schedule:  
 $r_1(A); A := A - 50; w_1(A); r_1(B); B := B + 50; w_1(B);$   
 $r_2(A); t := 0.1A; A := A - t; w_2(A); r_2(B); B := B + t; w_2(B);$
- Another schedule:  
 $r_1(A); A := A - 50; w_1(A); r_2(A); t := 0.1A; A := A - t;$   
 $w_2(A); r_1(B); B := B + 50; w_1(B); r_2(B); B := B + t; w_2(B);$
- This is not a serial schedule but is **equivalent**
- Yet another schedule:  
 $r_1(A); A := A - 50; r_2(A); t := 0.1A; A := A - t; w_2(A);$   
 $w_1(A); r_1(B); B := B + 50; w_1(B); r_2(B); B := B + t; w_2(B);$



# Example

- T1 transfers 50 from A to B and then T2 transfers 10% of A to B
- A **serial** schedule:  
 $r_1(A); A := A - 50; w_1(A); r_1(B); B := B + 50; w_1(B);$   
 $r_2(A); t := 0.1A; A := A - t; w_2(A); r_2(B); B := B + t; w_2(B);$
- Another schedule:  
 $r_1(A); A := A - 50; w_1(A); r_2(A); t := 0.1A; A := A - t;$   
 $w_2(A); r_1(B); B := B + 50; w_1(B); r_2(B); B := B + t; w_2(B);$
- This is not a serial schedule but is **equivalent**
- Yet another schedule:  
 $r_1(A); A := A - 50; r_2(A); t := 0.1A; A := A - t; w_2(A);$   
 $w_1(A); r_1(B); B := B + 50; w_1(B); r_2(B); B := B + t; w_2(B);$
- This is not a serial schedule and is not equivalent either

# Serializability

- Each transaction preserves database consistency
- Hence, a serial schedule also does that

# Serializability

- Each transaction preserves database consistency
- Hence, a serial schedule also does that
- A schedule is **serializable** if it is *equivalent* to a serial schedule
- There are different notions of equivalence
  - **Conflict serializability**
  - **View serializability**
- Operations other than *read* and *write* are ignored

# Serializability

- Each transaction preserves database consistency
- Hence, a serial schedule also does that
- A schedule is **serializable** if it is *equivalent* to a serial schedule
- There are different notions of equivalence
  - **Conflict serializability**
  - **View serializability**
- Operations other than *read* and *write* are ignored
- Instruction  $I_i$  of transaction  $T_i$  **conflicts** with  $I_j$  of  $T_j$  if and only if they access the **same** data item and at least one of them is a **write**

# Serializability

- Each transaction preserves database consistency
- Hence, a serial schedule also does that
- A schedule is **serializable** if it is *equivalent* to a serial schedule
- There are different notions of equivalence
  - **Conflict serializability**
  - **View serializability**
- Operations other than *read* and *write* are ignored
- Instruction  $I_i$  of transaction  $T_i$  **conflicts** with  $I_j$  of  $T_j$  if and only if they access the **same** data item and at least one of them is a **write**
- Intuitively, a conflict enforces a logical temporal order of the instructions
- Consequently, if two instructions do not conflict, they can be interchanged

# Conflict Serializability

- A schedule  $S$  is **conflict equivalent** to another schedule  $S'$  if it can be transformed to  $S'$  by a series of swaps of *non-conflicting* instructions

# Conflict Serializability

- A schedule  $S$  is **conflict equivalent** to another schedule  $S'$  if it can be transformed to  $S'$  by a series of swaps of *non-conflicting* instructions
- A schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule

# Conflict Serializability

- A schedule  $S$  is **conflict equivalent** to another schedule  $S'$  if it can be transformed to  $S'$  by a series of swaps of *non-conflicting* instructions
- A schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule
- A serial schedule is conflict serializable, but not vice versa
- If a schedule is conflict serializable, it is correct in the sense that it preserves database consistency



# Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$

# Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$   
is *conflict serializable* as it is conflict equivalent to the serial schedule  
 $T_1 T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$ 
  - It is not required to be conflict equivalent to  $T_2 T_1$  as well

# Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$   
is *conflict serializable* as it is conflict equivalent to the serial schedule  
 $T_1 T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$ 
  - It is not required to be conflict equivalent to  $T_2 T_1$  as well
- $r_1(a)w_2(a)w_1(a)$

# Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$   
is *conflict serializable* as it is conflict equivalent to the serial schedule  
 $T_1 T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$ 
  - It is not required to be conflict equivalent to  $T_2 T_1$  as well
- $r_1(a)w_2(a)w_1(a)$   
is *not* conflict serializable as it is not conflict equivalent to either of the two serial schedules  $T_1 T_2$  and  $T_2 T_1$

# View Serializability

- Two schedules are **view equivalent** if the reads in them get the same “**view**”, i.e., they read the value produced by the same write operation

# View Serializability

- Two schedules are **view equivalent** if the reads in them get the same “**view**”, i.e., they read the value produced by the same write operation
- Formally, two schedules  $S$  and  $S'$  are **view equivalent** if
  - 1 For each data item  $x$ , if a transaction  $T$  reads the initial value of  $x$  in  $S$ , it reads the same initial value of  $x$  in  $S'$  as well
  - 2 For each data item  $x$ , if a transaction  $T$  writes the final value of  $x$  in  $S$ , it writes the final value of  $x$  in  $S'$  as well
  - 3 If transaction  $T_i$  reads the value of data item  $x$  produced by write by transaction  $T_j$  in  $S$ , it must read the value written by  $T_j$  in  $S'$  as well
- A schedule  $S$  is **view serializable** if it is view equivalent to a serial schedule

# Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$

# Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$   
is *view serializable* as it is view equivalent to the serial schedule  
 $T_1 T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$



# Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$   
is *view serializable* as it is view equivalent to the serial schedule  
 $T_1 T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$
- $r_1(a)w_2(a)w_1(a)w_3(a)$

# Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$   
is *view serializable* as it is view equivalent to the serial schedule  
 $T_1 T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$
- $r_1(a)w_2(a)w_1(a)w_3(a)$   
is *view serializable* as it is view equivalent to the serial schedule  
 $T_1 T_2 T_3 : r_1(a)w_1(a)w_2(a)w_3(a)$

# Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$   
is *view serializable* as it is view equivalent to the serial schedule  
 $T_1 T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$
- $r_1(a)w_2(a)w_1(a)w_3(a)$   
is *view serializable* as it is view equivalent to the serial schedule  
 $T_1 T_2 T_3 : r_1(a)w_1(a)w_2(a)w_3(a)$
- $r_1(a)w_2(a)w_1(a)$

# Example

- $S : r_1(a)w_1(a)r_2(a)w_2(a)r_1(b)w_1(b)r_2(b)w_2(b)$   
is *view serializable* as it is view equivalent to the serial schedule  $T_1 T_2 : r_1(a)w_1(a)r_1(b)w_1(b)r_2(a)w_2(a)r_2(b)w_2(b)$
- $r_1(a)w_2(a)w_1(a)w_3(a)$   
is *view serializable* as it is view equivalent to the serial schedule  $T_1 T_2 T_3 : r_1(a)w_1(a)w_2(a)w_3(a)$
- $r_1(a)w_2(a)w_1(a)$   
is *not* view serializable as it is not view equivalent to either of the two serial schedules  $T_1 T_2$  and  $T_2 T_1$

# Relationship between View and Conflict Serializability

# Relationship between View and Conflict Serializability

- Every conflict serializable schedule is view serializable, but not vice versa
- Conflict serializability is *stricter* than view serializability

# Relationship between View and Conflict Serializability

- Every conflict serializable schedule is view serializable, but not vice versa
- Conflict serializability is *stricter* than view serializability
- They are same under the **constrained write assumption**
- In this assumption, every write of a data item  $x$  is constrained by the value of  $x$  it has read
  - $\text{write}(f(\text{read}(x)))$

# Relationship between View and Conflict Serializability

- Every conflict serializable schedule is view serializable, but not vice versa
- Conflict serializability is *stricter* than view serializability
- They are same under the **constrained write assumption**
- In this assumption, every write of a data item  $x$  is constrained by the value of  $x$  it has read
  - $\text{write}(f(\text{read}(x)))$
- With unconstrained writes (**blind writes**), a schedule that is view serializable is not necessarily conflict serializable
  - Blind writes: writes to a data item without reading it



# Relationship between View and Conflict Serializability

- Every conflict serializable schedule is view serializable, but not vice versa
- Conflict serializability is *stricter* than view serializability
- They are same under the **constrained write assumption**
- In this assumption, every write of a data item  $x$  is constrained by the value of  $x$  it has read
  - $\text{write}(f(\text{read}(x)))$
- With unconstrained writes (**blind writes**), a schedule that is view serializable is not necessarily conflict serializable
  - Blind writes: writes to a data item without reading it
- Every view serializable schedule that is not conflict serializable must have blind writes

# Other Notions of Equivalence

- Conflict and view serializable schedules are restrictive in the sense that they aim to guarantee database consistency without analyzing the result

# Other Notions of Equivalence

- Conflict and view serializable schedules are restrictive in the sense that they aim to guarantee database consistency without analyzing the result
- A schedule  $S$  is **result equivalent** to a schedule  $S'$  if it produces the same result as  $S'$
- Consider  
 $r_1(A); A := A - 50; w_1(A); r_2(B); B := B - 10; w_2(B);$   
 $r_1(B); B := B + 50; w_1(B); r_2(A); A := A + 10; w_2(A);$
- It produces the **same result** as the serial schedule  
 $r_1(A); A := A - 50; w_1(A); r_1(B); B := B + 50; w_1(B);$   
 $r_2(B); B := B - 10; w_2(B); r_2(A); A := A + 10; w_2(A);$   
but is not conflict or view serializable
- Determining such equivalence requires *semantic* analysis of operations other than read and write

# Testing for Serializability

- Create a **precedence graph** for the schedule
- Directed graph where each transaction is a vertex
- An edge from transaction  $T_i$  to  $T_j$  exists if
  - $w_i(x)$  precedes  $r_j(x)$ , or
  - $r_i(x)$  precedes  $w_j(x)$ , or
  - $w_i(x)$  precedes  $w_j(x)$

# Testing for Serializability

- Create a **precedence graph** for the schedule
- Directed graph where each transaction is a vertex
- An edge from transaction  $T_i$  to  $T_j$  exists if
  - $w_i(x)$  precedes  $r_j(x)$ , or
  - $r_i(x)$  precedes  $w_j(x)$ , or
  - $w_i(x)$  precedes  $w_j(x)$
- A schedule is conflict serializable if and only if its precedence graph is *acyclic*
- *Depth-first search* can detect cycles in  $O(n + m)$  time
- *Topological sorting* produces an equivalent serial order

# Testing for Serializability

- Create a **precedence graph** for the schedule
- Directed graph where each transaction is a vertex
- An edge from transaction  $T_i$  to  $T_j$  exists if
  - $w_i(x)$  precedes  $r_j(x)$ , or
  - $r_i(x)$  precedes  $w_j(x)$ , or
  - $w_i(x)$  precedes  $w_j(x)$
- A schedule is conflict serializable if and only if its precedence graph is *acyclic*
- *Depth-first search* can detect cycles in  $O(n + m)$  time
- *Topological sorting* produces an equivalent serial order
- Testing for view serializability is *NP-complete*
- Practical algorithms
  - Catches all non view serializable schedules
  - But can miss a view serializable schedule

# Recoverable Schedule

- Conflict and view serializability do not address failures
- Order of commits and aborts are important

# Recoverable Schedule

- Conflict and view serializability do not address failures
- Order of commits and aborts are important
- A schedule is called a **recoverable schedule** if
  - A transaction  $T_i$  reads a data item previously written by  $T_j$ , and
  - $T_j$  commits *before*  $T_i$  commits



# Recoverable Schedule

- Conflict and view serializability do not address failures
- Order of commits and aborts are important
- A schedule is called a **recoverable schedule** if
  - A transaction  $T_i$  reads a data item previously written by  $T_j$ , and
  - $T_j$  commits *before*  $T_i$  commits
- Consider  $r_1(a)w_1(a)r_2(a)r_1(b)$

# Recoverable Schedule

- Conflict and view serializability do not address failures
- Order of commits and aborts are important
- A schedule is called a **recoverable schedule** if
  - A transaction  $T_i$  reads a data item previously written by  $T_j$ , and
  - $T_j$  commits *before*  $T_i$  commits
- Consider  $r_1(a)w_1(a)r_2(a)r_1(b)$
- If  $T_2$  commits just after  $r_2(a)$ , i.e., if the schedule is  $r_1(a)w_1(a)r_2(a)c_2r_1(b)a_1$ , then it is *not* recoverable
  - If  $T_1$  crashes, then  $w_1(a)$  is undone, but  $T_2$  has already read a wrong value of  $a$  and committed

# Recoverable Schedule

- Conflict and view serializability do not address failures
- Order of commits and aborts are important
- A schedule is called a **recoverable schedule** if
  - A transaction  $T_i$  reads a data item previously written by  $T_j$ , and
  - $T_j$  commits *before*  $T_i$  commits
- Consider  $r_1(a)w_1(a)r_2(a)r_1(b)$
- If  $T_2$  commits just after  $r_2(a)$ , i.e., if the schedule is  $r_1(a)w_1(a)r_2(a)c_2r_1(b)a_1$ , then it is *not* recoverable
  - If  $T_1$  crashes, then  $w_1(a)$  is undone, but  $T_2$  has already read a wrong value of  $a$  and committed
- Therefore, to make it recoverable, the schedule should be  $r_1(a)w_1(a)r_2(a)r_1(b)c_1c_2$ 
  - If  $T_1$  aborts,  $T_2$  can also abort

# Cascading Rollbacks

- In recoverable schedules, a single transaction failure may lead to a series of rollbacks
- This is called **cascading rollbacks** or **cascading aborts**
- Consider  $r_1(a)w_1(a)r_2(a)w_2(a)r_3(a)r_1(b)a_1c_2c_3$
- It is recoverable
- However, if  $T_1$  fails,  $T_2$  and  $T_3$  must abort as well
- Not preferable as lot of work is undone

# Cascadeless Schedule

- Cascading rollbacks are eliminated

# Cascadeless Schedule

- Cascading rollbacks are eliminated
- A schedule is called a **cascadeless schedule** if
  - A transaction  $T_i$  reads a data item previously written by  $T_j$ , and
  - $T_j$  commits before  $T_i$  reads

# Cascadeless Schedule

- Cascading rollbacks are eliminated
- A schedule is called a **cascadeless schedule** if
  - A transaction  $T_i$  reads a data item previously written by  $T_j$ , and
  - $T_j$  commits before  $T_i$  reads
- Consider  $r_1(a)w_1(a)r_2(a)r_1(b)a_1c_2$

# Cascadeless Schedule

- Cascading rollbacks are eliminated
- A schedule is called a **cascadeless schedule** if
  - A transaction  $T_i$  reads a data item previously written by  $T_j$ , and
  - $T_j$  commits before  $T_i$  reads
- Consider  $r_1(a)w_1(a)r_2(a)r_1(b)a_1c_2$
- It is not cascadeless as  $T_2$  reads  $a$  written by  $T_1$  before  $T_1$  commits



# Cascadeless Schedule

- Cascading rollbacks are eliminated
- A schedule is called a **cascadeless schedule** if
  - A transaction  $T_i$  reads a data item previously written by  $T_j$ , and
  - $T_j$  commits before  $T_i$  reads
- Consider  $r_1(a)w_1(a)r_2(a)r_1(b)a_1c_2$
- It is not cascadeless as  $T_2$  reads  $a$  written by  $T_1$  before  $T_1$  commits
- Therefore, to make it cascadeless, the schedule should be  $r_1(a)w_1(a)r_1(b)c_1r_2(a)c_2$
- No completed transaction needs to be rolled back

# Cascadeless Schedule

- Cascading rollbacks are eliminated
- A schedule is called a **cascadeless schedule** if
  - A transaction  $T_i$  reads a data item previously written by  $T_j$ , and
  - $T_j$  commits before  $T_i$  reads
- Consider  $r_1(a)w_1(a)r_2(a)r_1(b)a_1c_2$
- It is not cascadeless as  $T_2$  reads  $a$  written by  $T_1$  before  $T_1$  commits
- Therefore, to make it cascadeless, the schedule should be  $r_1(a)w_1(a)r_1(b)c_1r_2(a)c_2$
- No completed transaction needs to be rolled back
- Every cascadeless schedule is recoverable, but not vice versa

# Strict Schedule

- Problem of writes remains in the sense that a later transaction may overwrite an uncommitted write

# Strict Schedule

- Problem of writes remains in the sense that a later transaction may overwrite an uncommitted write
- A schedule is called a **strict schedule** if
  - A transaction  $T_i$  reads or writes a data item previously written by  $T_j$ , and
  - $T_j$  commits before  $T_i$  *reads or writes*

# Strict Schedule

- Problem of writes remains in the sense that a later transaction may overwrite an uncommitted write
- A schedule is called a **strict schedule** if
  - A transaction  $T_i$  reads or writes a data item previously written by  $T_j$ , and
  - $T_j$  commits before  $T_i$  reads or writes
- Consider  $r_1(a)w_1(a)w_2(a)r_1(b)a_1c_2$

# Strict Schedule

- Problem of writes remains in the sense that a later transaction may overwrite an uncommitted write
- A schedule is called a **strict schedule** if
  - A transaction  $T_i$  reads or writes a data item previously written by  $T_j$ , and
  - $T_j$  commits before  $T_i$  reads or writes
- Consider  $r_1(a)w_1(a)w_2(a)r_1(b)a_1c_2$
- It is not strict as  $T_2$  writes  $a$  written by  $T_1$  before  $T_1$  commits

# Strict Schedule

- Problem of writes remains in the sense that a later transaction may overwrite an uncommitted write
- A schedule is called a **strict schedule** if
  - A transaction  $T_i$  reads or writes a data item previously written by  $T_j$ , and
  - $T_j$  commits before  $T_i$  reads or writes
- Consider  $r_1(a)w_1(a)w_2(a)r_1(b)a_1c_2$
- It is not strict as  $T_2$  writes  $a$  written by  $T_1$  before  $T_1$  commits
- Therefore, to make it strict, the schedule should be  $r_1(a)w_1(a)r_1(b)c_1w_2(a)c_2$

# Strict Schedule

- Problem of writes remains in the sense that a later transaction may overwrite an uncommitted write
- A schedule is called a **strict schedule** if
  - A transaction  $T_i$  reads or writes a data item previously written by  $T_j$ , and
  - $T_j$  commits before  $T_i$  reads or writes
- Consider  $r_1(a)w_1(a)w_2(a)r_1(b)a_1c_2$
- It is not strict as  $T_2$  writes  $a$  written by  $T_1$  before  $T_1$  commits
- Therefore, to make it strict, the schedule should be  $r_1(a)w_1(a)r_1(b)c_1w_2(a)c_2$
- Every strict schedule is cascadeless, but not vice versa



# Relationship among Schedules

