

Learning to Recommend via Matrix Factorization/Completion

Piyush Rai

Introduction to Machine Learning (CS771A)

October 30, 2018



Recommendation Systems

- The goal is to recommend more relevant items to users, based on previous interactions
- Used by many services



- Note: Relevance is subjective here
- A notion of relevance: I will buy/watch/like items that are similar to ones I did in the past
- Has been a very active research topic (in ML and allied areas) for a long time
 - Even a dedicated conference focusing on this topic specifically - RecSys



Recommendation Systems as Matrix Completion

- One of the most popular ways to solve the RecSys problem
- Suppose we have this partially complete ratings matrix

	Game of Thrones	Star Wars	Iron Man	Thor	Avengers	Guardians of the Galaxy
Male User 1	?	?	3	?	4	?
Female User 1	?	?	?	3	?	?
Male User 2	?	2	?	?	?	4
Female User 2	3	?	?	?	3	?
Male User 3	?	?	4	?	?	?
Female User 3	4	4	?	?	?	3

- Once completed, the completed matrix can be used to recommend “best” items for a given user
 - For example: Recommend the items that have a high (predicted) rating for the user
- Note: In addition to the user-item matrix, we may have **additional info** about the user/items
 - Some examples: User meta-data, item content description, user-user network, item-item similarity, etc.

Some Notation

	CUJI	Avatar	Iron Man	Iron Man	Spiderman	Lord of the Rings
Male User	?	?	3	?	4	?
Female User	?	?	?	3	?	?
User 3	?	2	?	?	?	4
User 4	3	?	?	?	3	?
User 5	?	?	4	?	?	?
User 6	4	4	?	?	?	3

- Let's denote the user-item $N \times M$ ratings matrix as \mathbf{X} (many entries missing)
- Suppose $\Omega = \{(n, m)\}$ is the set of indices for observed ratings
- Suppose Ω_{r_n} is the set of indices of items already rated by user n
- Suppose Ω_{c_m} is the set of indices of users who already rated item m

A Simple Heuristic: Item based “Collaborative Filtering”

	Game 1	Game 2	Game 3	Game 4	Game 5	Game 6
User 1	?	?	3	?	4	?
User 2	?	?	?	3	?	?
User 3	?	2	?	?	?	4
User 4	3	?	?	?	3	?
User 5	?	?	4	?	?	?
User 6	4	4	?	?	?	3

- For each user-item pair (n, m) , compute the missing rating X_{nm} as

$$X_{nm} \approx \frac{1}{|\Omega_{r_n}|} \sum_{m' \in \Omega_{r_n}} S_{mm'}^{(I)} X_{nm'}$$

where $S_{mm'}^{(I)} \in (0, 1)$ is the similarity between items m and m' (suppose known)

A Simple Heuristic: User based “Collaborative Filtering”

	Game 1	Game 2	Game 3	Game 4	Game 5	Game 6
User 1	?	?	3	?	4	?
User 2	?	?	?	3	?	?
User 3	?	2	?	?	?	4
User 4	3	?	?	?	3	?
User 5	?	?	4	?	?	?
User 6	4	4	?	?	?	3

- For each user-item pair (n, m) , compute the missing rating X_{nm} as

$$X_{nm} \approx \frac{1}{|\Omega_{c_m}|} \sum_{n' \in \Omega_{c_m}} S_{nn'}^{(U)} X_{n'm}$$

where $S_{nn'}^{(U)} \in (0, 1)$ is the similarity between users n and n' (suppose known)

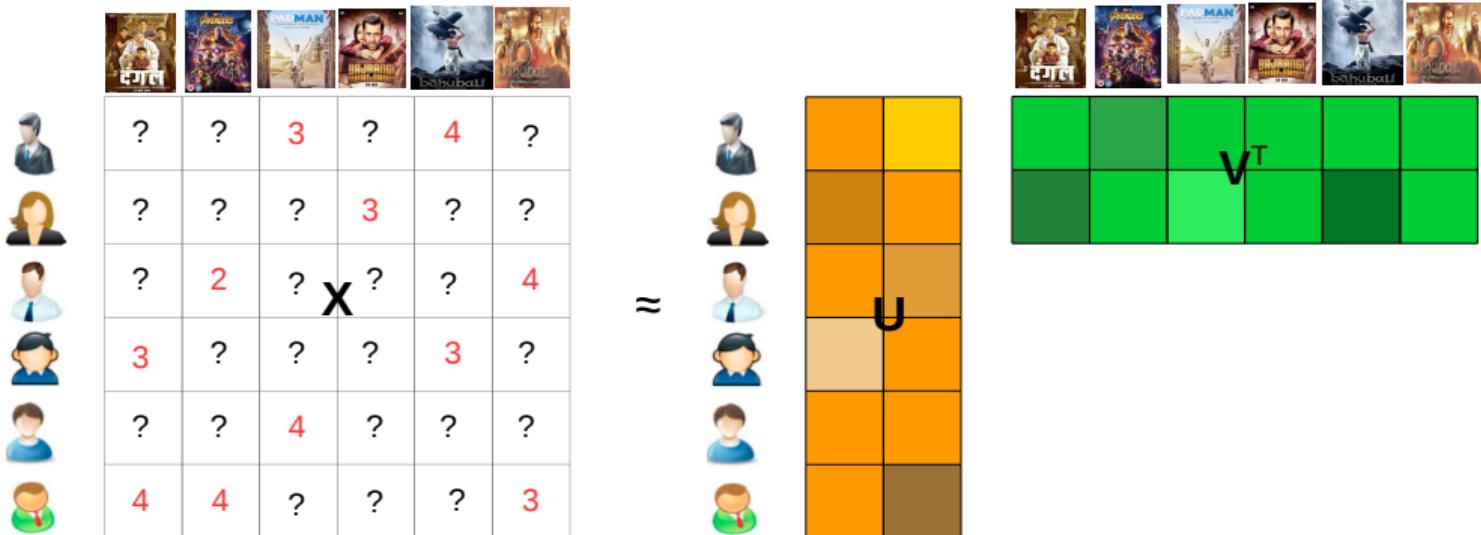
Limitations of Item/User Based Approach

	EDEN	LOST	MAN	IRON MAN	WING	UNIVERSE
Police Officer	?	?	3	?	4	?
Businesswoman	?	?	?	3	?	?
Male	?	2	?	?	?	4
Female	3	?	?	?	3	?
Teenager	?	?	4	?	?	?
Child	4	4	?	?	?	3

- User-User or Item-Item similarities may not be known beforehand
- We may have very little data in the user-item matrix and averaging may not be reliable



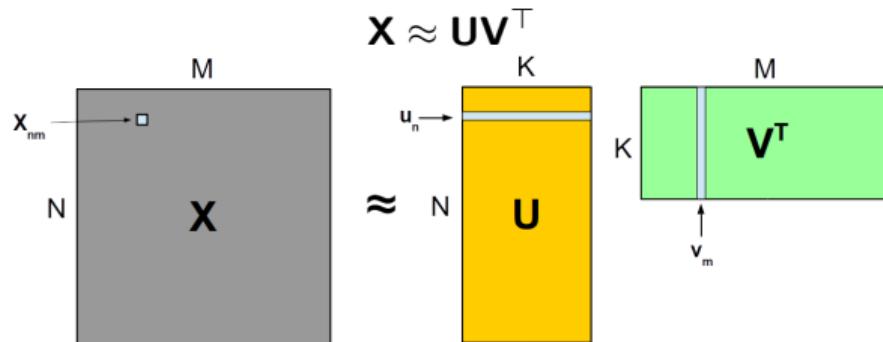
Towards a better approach: Matrix Factorization



If we can do the above factorization then any missing $X_{nm} \approx u_n^\top v_m$

Matrix Factorization

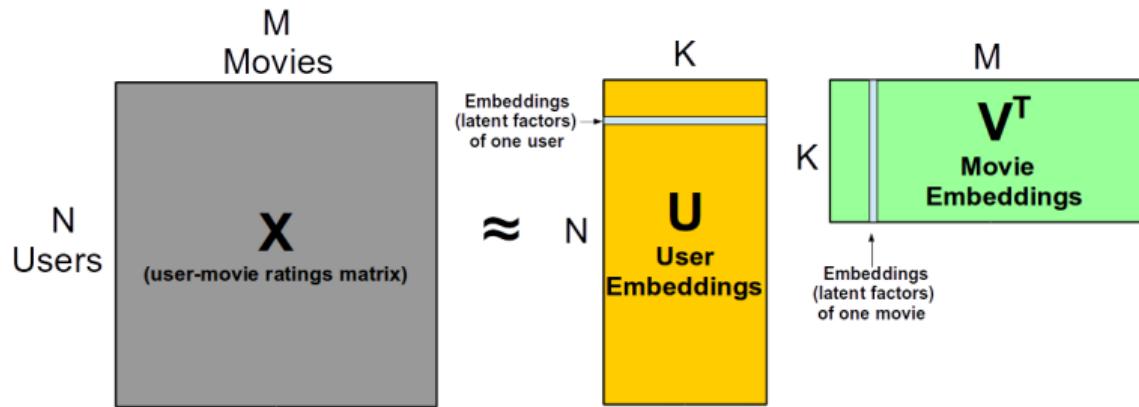
- Given a matrix \mathbf{X} of size $N \times M$, approximate it as a product of two matrices



- \mathbf{U} : $N \times K$ latent factor matrix
 - Each row of \mathbf{U} represents a K -dim latent factor u_n
- \mathbf{V} : $M \times K$ latent factor matrix
 - Each row of \mathbf{V} represents a K -dim latent factor v_n
 - Each entry of \mathbf{X} can be written as: $X_{nm} \approx u_n^T v_m = \sum_{k=1}^K u_{nk} v_{mk}$

Why Matrix Factorization?

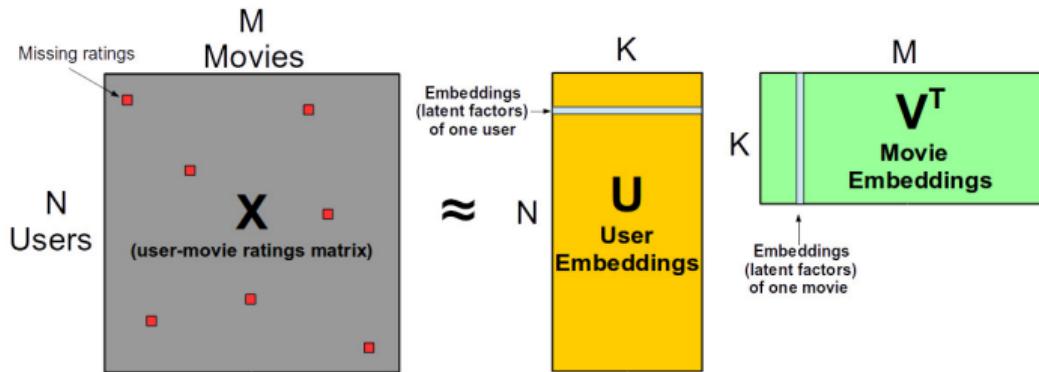
- The latent factors can be used/interpreted as “embeddings” or “learned features”



- Especially useful for learning good features for “dyadic” or relational data
 - Examples: Users-Movies ratings, Users-Products purchases, etc.
- If $K \ll \min\{M, N\}$ ⇒ then can also be seen as dimensionality reduction or a “low-rank factorization” of the matrix \mathbf{X} (somewhat like SVD)

Why Matrix Factorization?

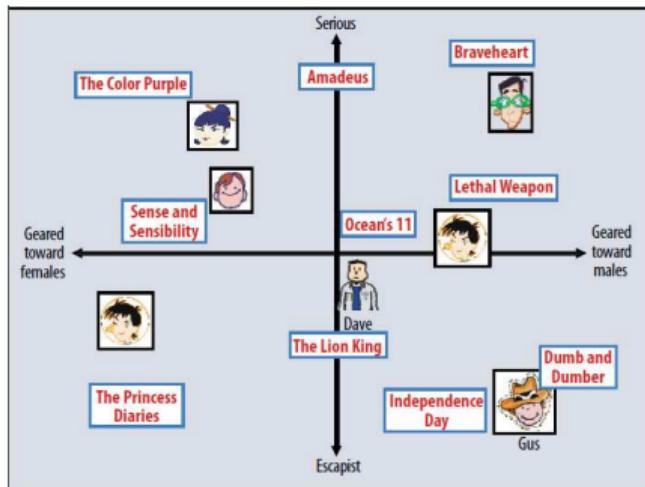
- Can also predict the missing/unknown entries in the original matrix



- Yes. \mathbf{U} and \mathbf{V} can be learned even when the matrix \mathbf{X} is only **partially observed** (we'll see shortly)
- After learning \mathbf{U} and \mathbf{V} , any missing X_{nm} can be approximated by $\mathbf{u}_n^\top \mathbf{v}_m$
- \mathbf{UV}^\top is the best low-rank matrix that approximates the full \mathbf{X}
- Note: The “[Netflix Challenge](#)” was won by a matrix factorization method

Interpreting the Embeddings/Latent Factors

- Embeddings/latent factors can often be interpreted. E.g., as “genres” if \mathbf{X} represents a user-movie rating matrix. A cartoon with $K = 2$ shown below

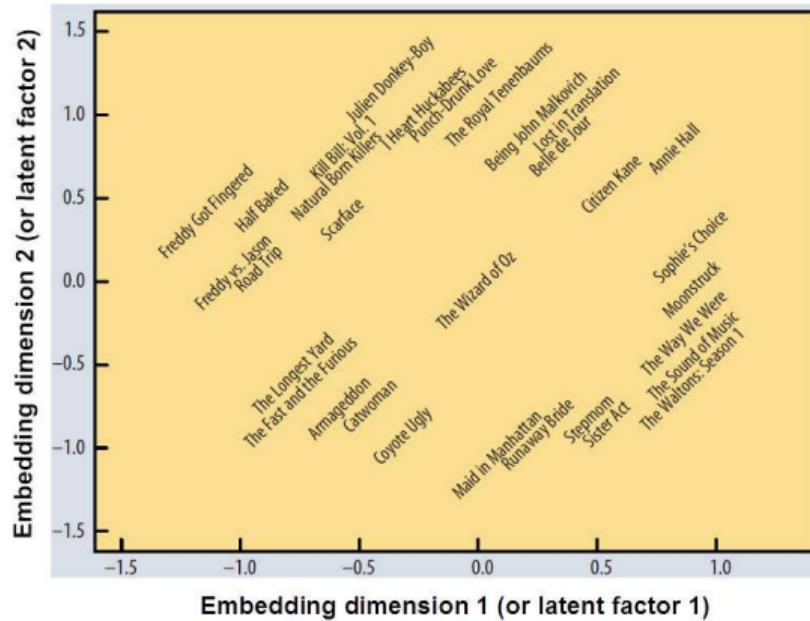


- Similar things (users/movies) get embedded nearby in the embedding space (two things will be deemed similar if their embeddings are similar). Thus useful for **computing similarities** and/or **making recommendations**

Picture courtesy: Matrix Factorization Techniques for Recommender Systems: Koren et al, 2009

Interpreting the Embeddings/Latent Factors

- Another illustration of 2-D embeddings of the movies only



- Similar movies will be embedded at nearby locations in the embedding space

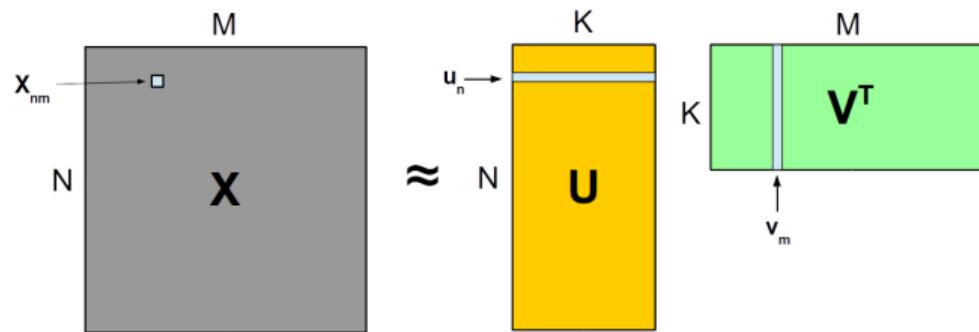
Picture courtesy: Matrix Factorization Techniques for Recommender Systems: Koren et al, 2009

Solving Matrix Factorization



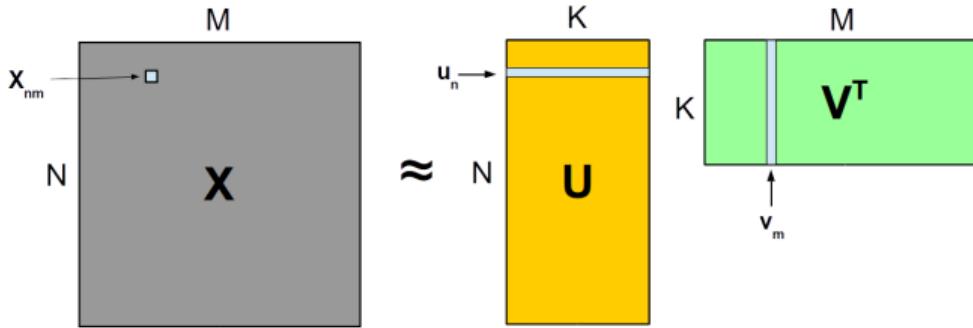
Matrix Factorization

- Recall our matrix factorization model: $\mathbf{X} \approx \mathbf{U}\mathbf{V}^\top$
- Goal: learn \mathbf{U} and \mathbf{V} , given a subset Ω of entries in \mathbf{X} (let's call it \mathbf{X}_Ω)
- Recall our notations:
 - $\Omega = \{(n, m)\}$: X_{nm} is observed
 - Ω_{r_n} : column indices of observed entries in row n of \mathbf{X}
 - Ω_{c_m} : row indices of observed entries in column m of \mathbf{X}



Matrix Factorization

- We want \mathbf{X} to be as close to \mathbf{UV}^\top as possible



- Let's define a squared "loss function" over the observed entries in \mathbf{X}

$$\mathcal{L} = \sum_{(n,m) \in \Omega} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2$$

- Here the latent factors $\{\mathbf{u}_n\}_{n=1}^N$ and $\{\mathbf{v}_m\}_{m=1}^M$ are the **unknown parameters**
- Squared loss chosen only for simplicity; other loss functions can be used
- How do we learn $\{\mathbf{u}_n\}_{n=1}^N$ and $\{\mathbf{v}_m\}_{m=1}^M$?



Alternating Optimization

- We will use an ℓ_2 regularized version of the squared loss function

$$\mathcal{L} = \sum_{(n,m) \in \Omega} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \sum_{n=1}^N \lambda_U \|\mathbf{u}_n\|^2 + \sum_{m=1}^M \lambda_V \|\mathbf{v}_m\|^2$$

- A **non-convex** problem. Difficult to optimize w.r.t. \mathbf{u}_n and \mathbf{v}_m jointly.
- One way is to solve for \mathbf{u}_n and \mathbf{v}_m in an **alternating fashion**, e.g.,

- $\forall n$, fix all variables except \mathbf{u}_n and solve the optim. problem w.r.t. \mathbf{u}_n

$$\arg \min_{\mathbf{u}_n} \sum_{m \in \Omega_{r_n}} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \lambda_U \|\mathbf{u}_n\|^2$$

- $\forall m$, fix all variables except \mathbf{v}_m and solve the optim. problem w.r.t. \mathbf{v}_m

$$\arg \min_{\mathbf{v}_m} \sum_{n \in \Omega_{c_m}} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \lambda_V \|\mathbf{v}_m\|^2$$

- Iterate until not converged
- Each of these subproblems has a simple, convex objective function
- Convergence properties of such methods have been studied extensively

The Solutions

- Easy to show that the problem

$$\arg \min_{\mathbf{u}_n} \sum_{m \in \Omega_{r_n}} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \lambda_U \|\mathbf{u}_n\|^2$$

.. has the solution

$$\mathbf{u}_n = \left(\sum_{m \in \Omega_{r_n}} \mathbf{v}_m \mathbf{v}_m^\top + \lambda_U \mathbf{I}_K \right)^{-1} \left(\sum_{m \in \Omega_{r_n}} X_{nm} \mathbf{v}_m \right)$$

- Likewise, the problem

$$\arg \min_{\mathbf{v}_m} \sum_{n \in \Omega_{c_m}} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \lambda_V \|\mathbf{v}_m\|^2$$

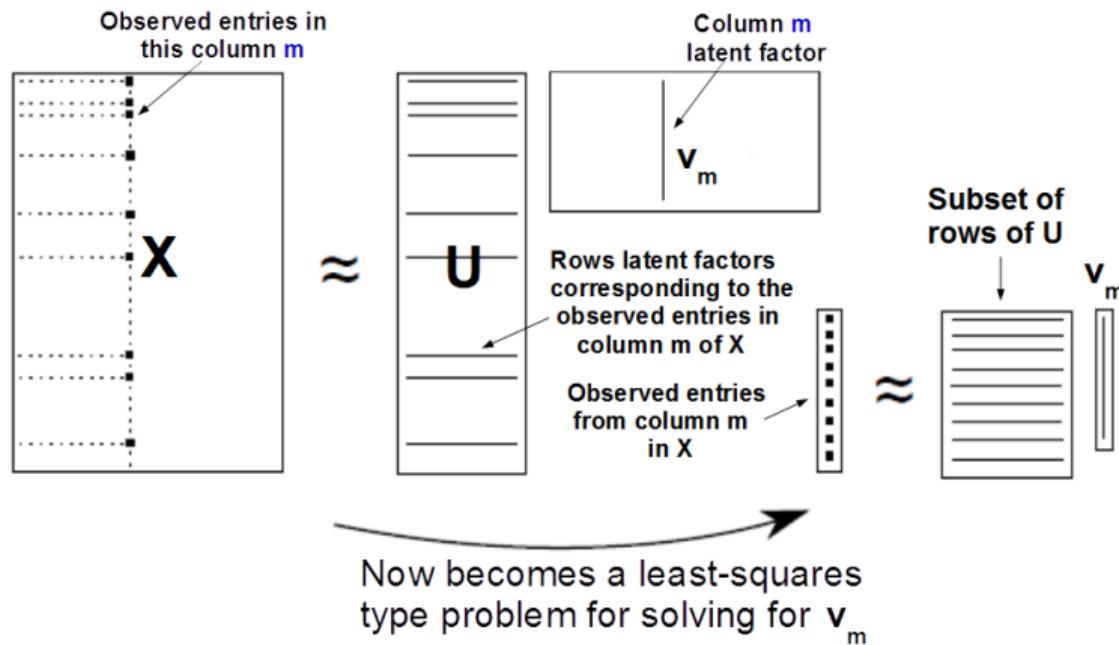
.. has the solution

$$\mathbf{v}_m = \left(\sum_{n \in \Omega_{c_m}} \mathbf{u}_n \mathbf{u}_n^\top + \lambda_V \mathbf{I}_K \right)^{-1} \left(\sum_{n \in \Omega_{c_m}} X_{nm} \mathbf{u}_n \right)$$

- Note that this is very similar to (regularized) least squares regression
- Thus matrix factorization can be also seen as a sequence of regression problems (one for each latent factor)

Matrix Factorization as Regression

Suppose we are solving for v_m (with U and all other v_m 's fixed)



Can think of solving for u_n (with V and all other u_n 's fixed) in the same way



Matrix Factorization as Regression

- A very useful way to understand matrix factorization
- Can modify the regularized least-squares like objective

$$\arg \min_{\mathbf{u}_n} \sum_{m \in \Omega_{r_n}} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \lambda_U \mathbf{u}_n^\top \mathbf{u}_n$$

.. using other loss functions and regularizers

- Some possible modifications:
 - If entries in the matrix \mathbf{X} are binary, counts, etc. then we can replace the squared loss function by some other loss function (e.g., logistic or Poisson)
 - Can impose other constraints on the latent factors, e.g., non-negativity, sparsity, etc. (by changing the regularizer)
 - Can think of this also as a probabilistic model (a likelihood function on X_{nm} and priors on the latent factors $\mathbf{u}_n, \mathbf{v}_m$) and do MLE/MAP

Matrix Factorization: The Complete Algorithm

- Input: Partially complete matrix \mathbf{X}_Ω
- Initialize the latent factors $\mathbf{v}_1, \dots, \mathbf{v}_M$ randomly
- Iterate until not converged
 - Update each row latent factor $\mathbf{u}_n, n = 1, \dots, N$ (can be in parallel)

$$\mathbf{u}_n = \left(\sum_{m \in \Omega_{r_n}} \mathbf{v}_m \mathbf{v}_m^\top + \lambda_U \mathbf{I}_K \right)^{-1} \left(\sum_{m \in \Omega_{r_n}} X_{nm} \mathbf{v}_m \right)$$

- Update each column latent factor $\mathbf{v}_m, m = 1, \dots, M$ (can be in parallel)
- Final prediction for any (missing) entry: $X_{nm} = \mathbf{u}_n^\top \mathbf{v}_m$



A Faster Algorithm via SGD

- Alternating optimization is nice but can be slow (note that it requires matrix inversion with cost $O(K^3)$ for updating each latent factor $\mathbf{u}_n, \mathbf{v}_m$)
- An alternative is to use stochastic gradient descent (SGD). In each round, select a randomly chosen entry X_{nm} with $(n, m) \in \Omega$
- Consider updating \mathbf{u}_n . For loss function $\sum_{m \in \Omega_{r_n}} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \lambda_U \|\mathbf{u}_n\|^2$, the stochastic gradient w.r.t. \mathbf{u}_n using this randomly chosen entry X_{nm} is

$$-(X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m) \mathbf{v}_m + \lambda_U \mathbf{u}_n$$

- Thus the SGD update for \mathbf{u}_n will be

$$\mathbf{u}_n = \mathbf{u}_n - \eta(\lambda_U \mathbf{u}_n - (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m) \mathbf{v}_m)$$

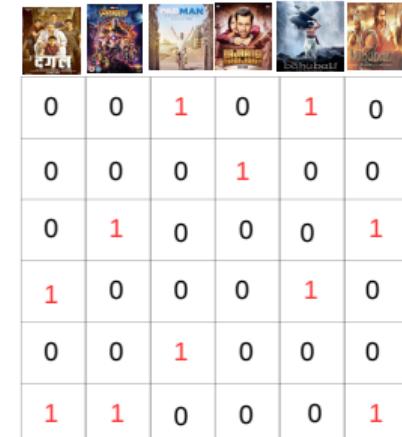
- Likewise, the SGD update for \mathbf{v}_m will be

$$\mathbf{v}_m = \mathbf{v}_m - \eta(\lambda_V \mathbf{v}_m - (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m) \mathbf{u}_n)$$

- The SGD algorithm chooses a random entry X_{nm} in each iteration, updates $\mathbf{u}_n, \mathbf{v}_m$, and repeats until convergence (\mathbf{u}_n 's, \mathbf{v}_m 's randomly initialized).

Explicit Feedback vs Implicit Feedback Data

- Often the user-item matrix \mathbf{X} is a binary matrix
- $X_{nm} = 1$ means user n watched and liked on the item m
- What does $X_{nm} = 0$ mean? Watched but didn't like?
- Or $X_{nm} = 0$ mean user wasn't exposed to this item?
- There is no way to distinguish such 0s in \mathbf{X}
- Such binary \mathbf{X} is called “**implicit feedback**” as opposed to explicit feedback (e.g., ratings)
- Such data needs more careful modeling (other loss functions, not squared/logistic)
- Some popular schemes include
 - Downweighting the contribution of 0s in the loss function
 - Use ranking based loss function, e.g., want $\mathbf{u}_n^\top \mathbf{v}_m > \mathbf{u}_n^\top \mathbf{v}_{m'}$ if $X_{nm} = 1$ and $X_{nm'} = 0$



	Item 1	Item 2	Item 3	Item 4	Item 5	Item 6
User 1	0	0	1	0	1	0
User 2	0	0	0	1	0	0
User 3	0	1	0	0	0	1
User 4	1	0	0	0	1	0
User 5	0	0	1	0	0	0
User 6	1	1	0	0	0	1

Inductive Matrix Completion

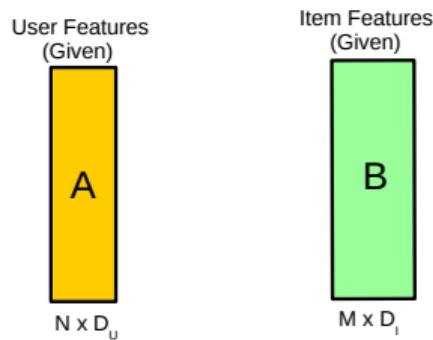
- “Inductive” here means that we would like to extrapolate to new users/items
- Also known as the “cold-start” problem in RecSys literature



- The matrix factorization approach would need latent factors for the new users/items
- How to compute these latent factors without any ratings for such users/items?

Inductive Matrix Completion

- Often we have some additional “meta-data” about the users or items (or both)
 - Example: User profile info, item description/image, etc.
- Can use this meta-data to get some features
- Assume we have D_u features for each user and D_I features for each item

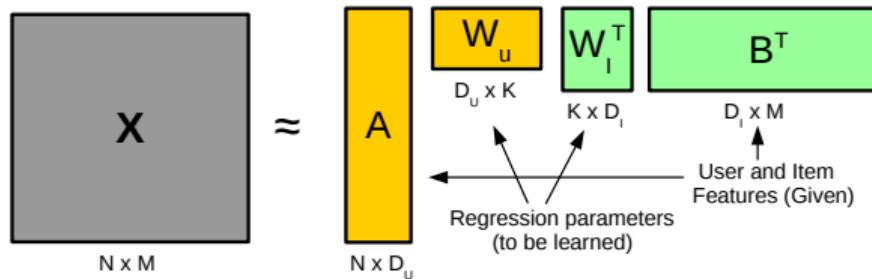


- One possibility now: Use these features/meta-data to get the latent factors for users/items

Matrix Factorization for Inductive Matrix Completion

- Basic idea: Assume $\mathbf{X} \approx \mathbf{U}\mathbf{V}^\top$ but regress \mathbf{U} and \mathbf{V} using \mathbf{A} and \mathbf{B} , respectively

$$\mathbf{U} = \mathbf{A}\mathbf{W}_u \quad \text{and} \quad \mathbf{V} = \mathbf{B}\mathbf{W}_I$$



- The loss function will be

$$\|\mathbf{X} - \mathbf{UV}^\top\|^2 = \|\mathbf{X} - (\mathbf{AW}_u) \times (\mathbf{BW}_I)^\top\|^2$$

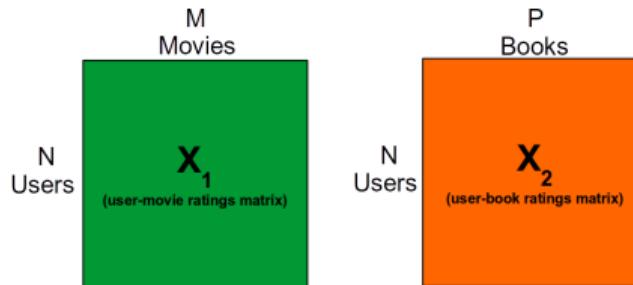
- We optimize this loss function w.r.t. \mathbf{W}_U and \mathbf{W}_I
- For a new user with features \mathbf{a}_* , we compute the latent factor $\mathbf{u}_* = \mathbf{a}_*\mathbf{W}_U$
- For a new item with features \mathbf{b}_* , we compute the latent factor $\mathbf{v}_* = \mathbf{b}_*\mathbf{W}_I$

Some Other Extensions of Matrix Factorization



Joint Matrix Factorization

- Can do joint matrix factorization of more than one matrices
- Consider two “ratings” matrices with the N users shared in both



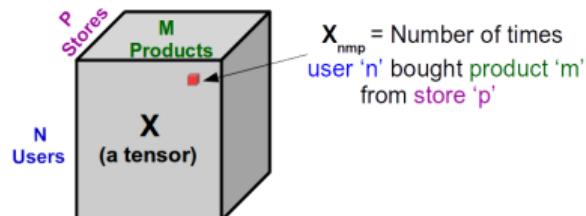
- Can assume the following matrix factorization

$$\mathbf{X}_1 \approx \mathbf{U} \mathbf{V}_1^\top \quad \text{and} \quad \mathbf{X}_2 \approx \mathbf{U} \mathbf{V}_2^\top$$

- Note that the user latent factor matrix \mathbf{U} is shared in both factorizations
- Gives a way to [learn features by combining multiple data sets](#) (2 in this case)
- Can use the alternating optimization to solve for \mathbf{U} , \mathbf{V}_1 and \mathbf{V}_2

Tensor Factorization

- A “tensor” is a generalization of a matrix to more than two dimensions
- Consider a 3-dim (or 3-mode or 3-way) tensor \mathbf{X} of size $N \times M \times P$



- We can model each entry of tensor \mathbf{X} as

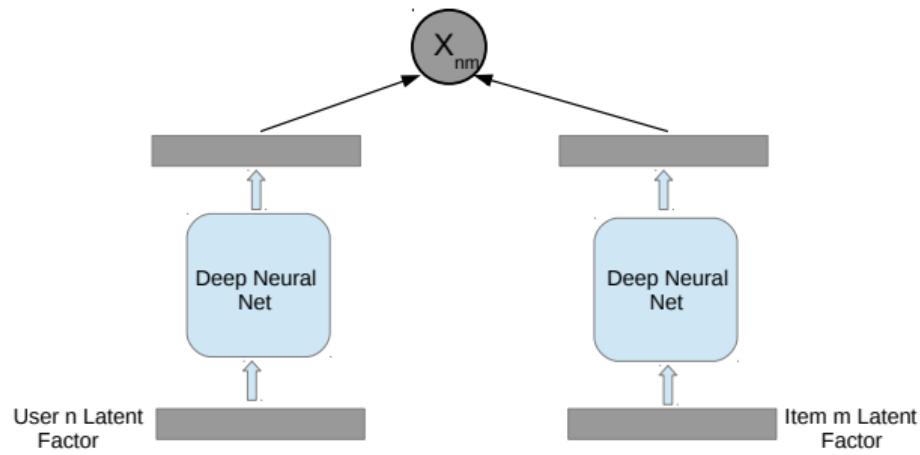
$$X_{nmp} \approx \mathbf{u}_n \odot \mathbf{v}_m \odot \mathbf{w}_p = \sum_{k=1}^K u_{nk} v_{mk} w_{pk}$$

- Can learn $\{\mathbf{u}_n\}_{n=1}^N, \{\mathbf{v}_m\}_{m=1}^M, \{\mathbf{w}_p\}_{p=1}^P$ using alternating optimization
- These K -dim. “embeddings” can be used as features for other tasks (e.g., tensor completion, computing similarities, etc.)
- The model also be easily extended to tensors having than 3 dimensions
- Several specialized algorithms for tensor factorization (CP/Tucker decomposition, etc.)



(And of course..) Deep Learning based Recommender Systems

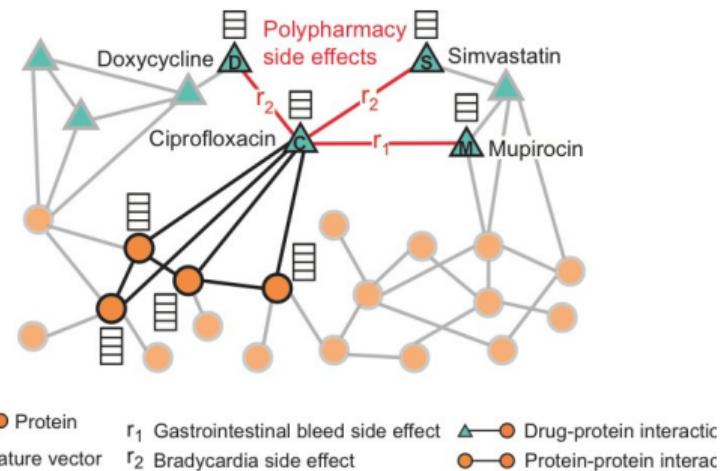
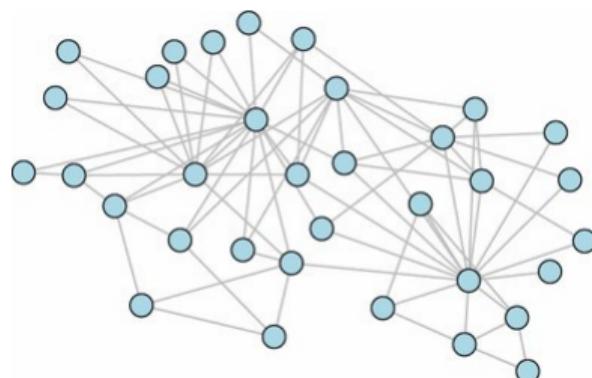
Basic idea: Matrix entries are nonlinear transformations of the latent factors $X_{nm} \approx f(\mathbf{u}_n)^T f(\mathbf{v}_m)$



This above is a simple version. Many more sophisticated variants exist (posted a reference on the course webpage in case you are interested in deep learning methods for recommender systems)

Applications to Link Prediction in Graphs

- The user-item matrix is like a bipartite graph
- The matrix factorization ideas we saw today can also be used for any type of graph



- Thus we can get node embeddings as well as a way to do link prediction in such graphs

Right side picture courtesy: snap.stanford.edu

Some Final Comments..

- Looked at some basic as well as some state-of-the-art approaches for recommendation systems
- Matrix factorization/completion is one of the dominant approaches (though not the only one) to solve the problem
- Other want to take into account other criteria such as **freshness and diversity** of recommendations
 - Don't want to keep recommending the similar items again and again
- Often helps to incorporate sources (e.g., meta data) other than just the user-item matrix
 - We saw some techniques already (e.g., inductive matrix completion)
- Temporal nature can also be incorporated (e.g., user and item latent factors may evolve in time)
- Still an ongoing area of active research

